

# Performance Evaluation

## AWS-EC2 Cloud Benchmark

This document contains the performance evaluation of CPU Benchmark, Disk Benchmark and Memory Benchmark on AWS EC-2 Cloud. My experiments cover all the original specification for all benchmarks of AWS-EC2 cloud. There are so many experiments done for all benchmarks and every experiment has been done 3-times for calculating average and standard deviation of all figures.

First of all, I describe the Environment of AWS-EC2 Cloud which is shown in below figure. Then the specification of each benchmark is presented followed by experiment analysis and result in form of table and graph.

### 1. AWS-EC2 Environment

write this command on AWS-EC2 instance after launching. It gives you following output of all specification

```
[ec2-user@ip-172-31-49-170 ~]$ lscpu
```

```
[ec2-user@ip-172-31-49-170 ~]$ clear
[ec2-user@ip-172-31-49-170 ~]$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                1
On-line CPU(s) list:   0
Thread(s) per core:    1
Core(s) per socket:    1
Socket(s):             1
NUMA node(s):         1
Vendor ID:             GenuineIntel
CPU family:            6
Model:                63
Model name:            Intel(R) Xeon(R) CPU E5-2676 v3 @ 2.40GHz
Stepping:              2
CPU MHz:               2400.046
BogoMIPS:              4800.09
Hypervisor vendor:     Xen
Virtualization type:   full
L1d cache:             32K
L1i cache:             32K
L2 cache:              256K
L3 cache:              30720K
NUMA node0 CPU(s):    0
[ec2-user@ip-172-31-49-170 ~]$
```

## 2. Benchmark Experiment and Result Analysis

I represent all experimental result analysis of three benchmark of CPU, Disk and Memory and explanations of each and every outcomes with comparison of theoretical performance.

### 2.1 CPU Benchmark

- For CPU Benchmark, calculating results in terms of FLOPS(FLoating point Operations per Second) and IOPS (Interger Operations per Second) and then find GFLOPS and GIOPS . (G – Giga  $10^9$ ).
- Results of FLOPS and IOPS according to no. of threads(1, 2 and 4 threads) shown in figure Fig 1. below in form of graph.

#### CPU Benchmark FLOPS and IOPS

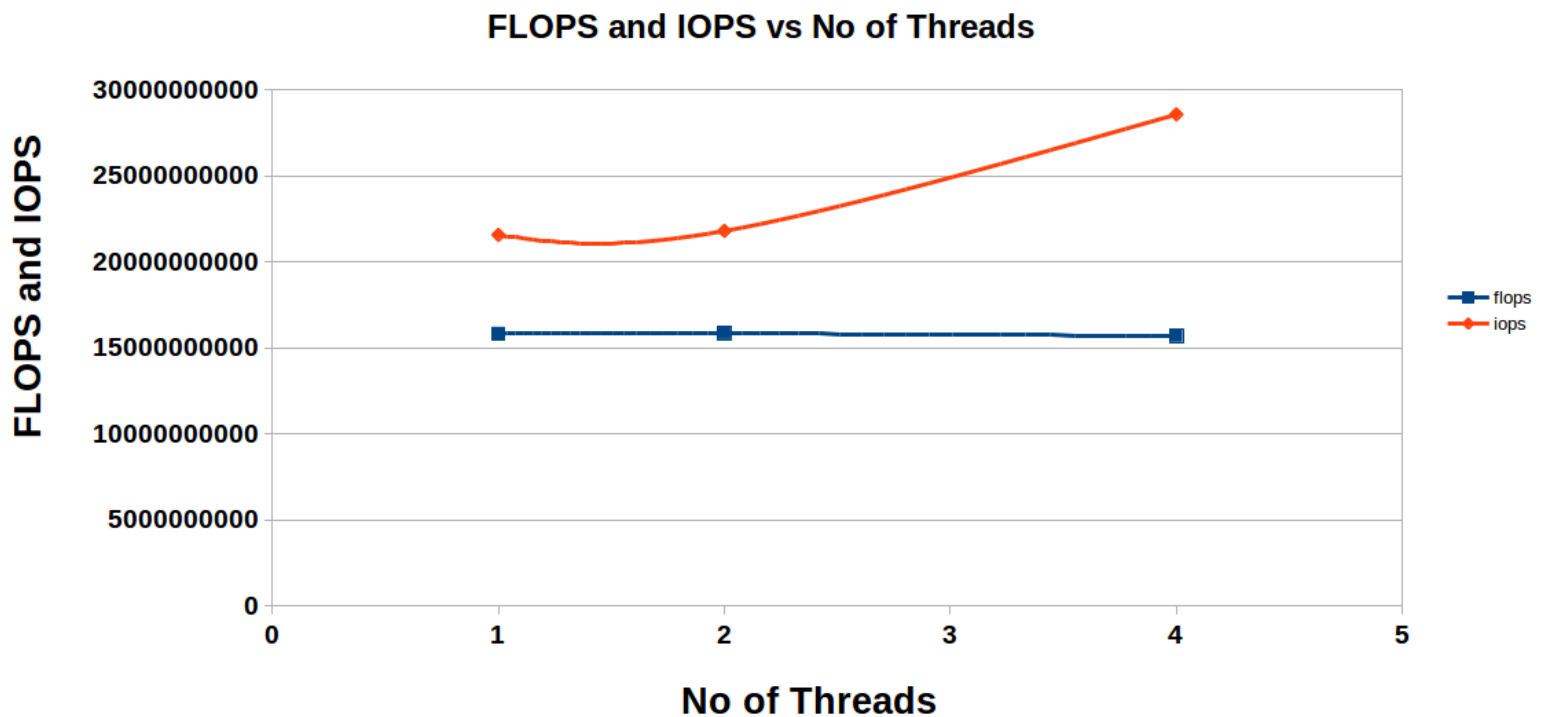


Fig 1.

c. According to practical result no. of optimal threads is “4” for best performance

No. of Threads	Average(GFLOPS)	SD(GFLOPS)	Average(GIOPS)	SD(GIOPS)
1	15.794472300	0.00934	21.573456800	0.17122
2	15.810272300	0.01747	21.642424300	0.11734
4	15.871454500	0.00794	28.533168500	0.12344

d. Practice performance

FLOPS = **15810272300** and GFLOPS = **15.81**  
 IOPS = **21563328500** and GIOPS = **21.63**

e. Theoretical peak performance of AWS-EC2 cloud is given below:

performance = (CPU Speed in GHz) \* (CPU IPC)\* (no.of cores in CPU) \* (number of CPU per node)

performance =  $2.4 * 16 * 1 * 1 = 38.4$  GFLOPS

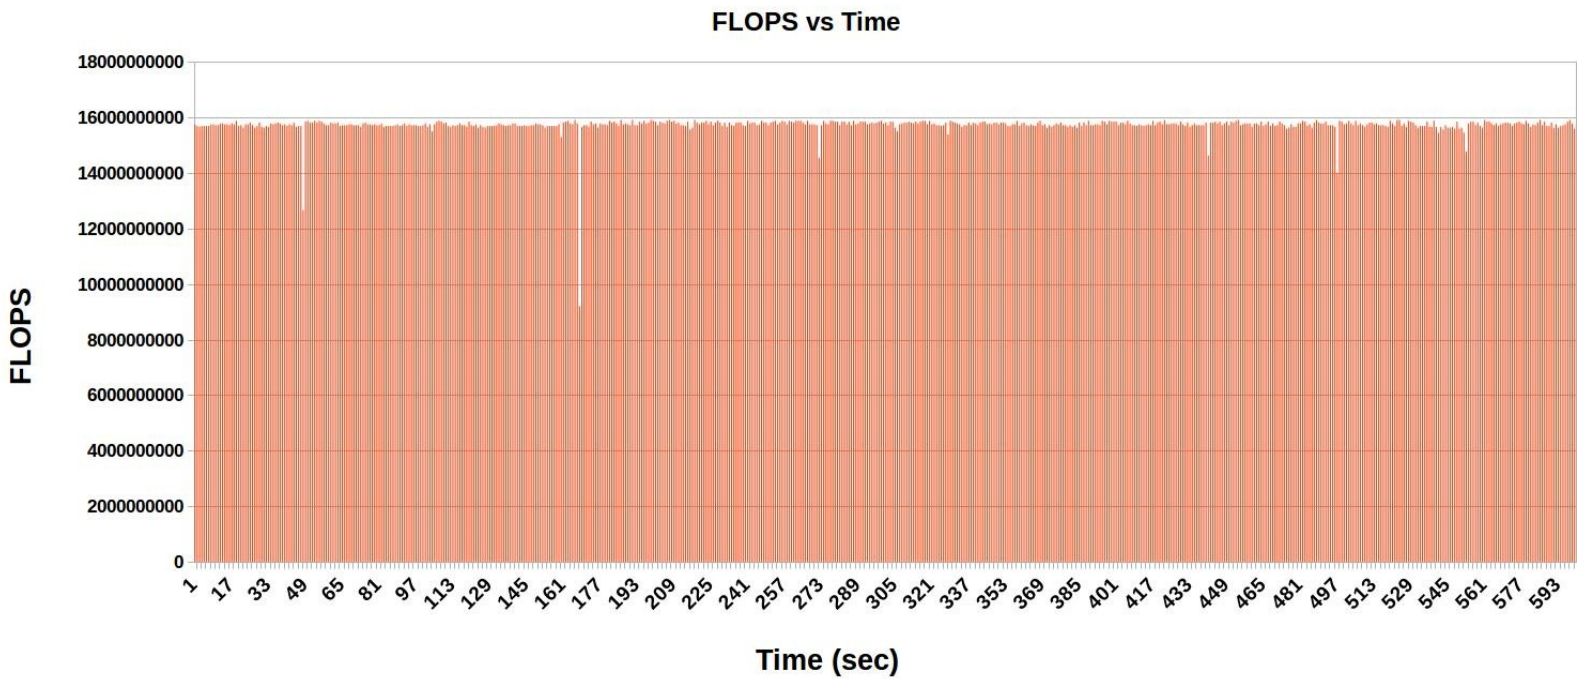
f. Efficiency of CPU Speed (GFLOPS) (**49.65 %**)

Efficiency =  $(15.81)/(38.4)$   
 = **49.65 %**

g. I run this program for 10 min and took 600 samples of 1 second each. I took samples of FLOPS and IOPS. It gives me a better idea of CPU performance while performing this much of experiments.


Given below graph gives better idea of FLOPS and IOPS

## Graph of FLOPS vs Time for 600 Samples



h. 10 mins samples of IOPS and measured CPU Performance.

## IOPS - 600 Samples (each per 1 second)



I. Practical Performance by running “linpack” on the AWS-EC2 Cloud. Below is the figure to see the performance of CPU benchmark.

$$\text{Efficiency of linpack result to theoretical result} = (33.69/38.4) = 87.73\%$$

Efficiency = 87.73%.

Figure is for “linpack” CPU Benchmark.

Fig 3.

```
[ec2-user@ip-172-31-49-170 linpack]$ ./runme_xeon64
This is a SAMPLE run script for SMP LINPACK. Change it to reflect
the correct number of CPUs/threads, problem input files, etc..
Fri Feb 12 00:19:36 UTC 2016
Intel(R) Optimized LINPACK Benchmark data

Current date/time: Fri Feb 12 00:19:36 2016

CPU frequency:      2.992 GHz
Number of CPUs: 1
Number of cores: 1
Number of threads: 1

Parameters are set to:

Number of tests: 15
Number of equations to solve (problem size) : 1000  2000  5000  10000 15000 18000 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array                  : 1000  2000  5008  10000 15000 18008 20016 22008 25000 26000 27000 30000 35000 40000 45000
Number of trials to run                     : 4      2      2      2      2      2      2      2      2      2      1      1      1      1      1
Data alignment value (in Kbytes)            : 4      4      4      4      4      4      4      4      4      4      4      1      1      1      1

Maximum memory requested that can be used=800204096, at the size=10000

===== Timing linear equation system solver =====

Size  LDA   Align. Time(s)   GFlops  Residual   Residual(norm) Check
1000  1000   4       0.026   26.1378  9.632295e-13 3.284860e-02 pass
1000  1000   4       0.025   26.2410  9.632295e-13 3.284860e-02 pass
1000  1000   4       0.025   26.4202  9.632295e-13 3.284860e-02 pass
1000  1000   4       0.026   26.0428  9.632295e-13 3.284860e-02 pass
2000  2000   4       0.195   27.3354  4.746648e-12 4.129002e-02 pass
2000  2000   4       0.193   27.7365  4.746648e-12 4.129002e-02 pass
5000  5008   4       2.494   33.4359  2.651185e-11 3.696863e-02 pass
5000  5008   4       2.489   33.5038  2.651185e-11 3.696863e-02 pass
10000 10000  4      18.682  35.6959  9.014595e-11 3.178637e-02 pass
10000 10000  4      18.440  36.1648  9.014595e-11 3.178637e-02 pass

Performance Summary (GFlops)

Size  LDA   Align.  Average  Maximal
1000  1000   4       26.2104  26.4202
2000  2000   4       27.5359  27.7365
5000  5008   4       33.4698  33.5038
10000 10000  4       35.9303  36.1648

Residual checks PASSED

End of tests

Done: Fri Feb 12 00:20:25 UTC 2016
[ec2-user@ip-172-31-49-170 linpack]$
```

## Explanations

CPU performance in terms of GFLOPS and GIOPS is gradually increase initially when doing millions of operations. Then after it is stable. So in the graph is shows stable value for billions of operations. We need to do billions of operations for CPU performance for better and consistent result. It gives accurate answer for 100000000 iterations for 1s and I took 600 samples by 10 mins so that I got consistent result.

## Issued Faced

While measuring the CPU benchmark first thing I m confused with the IPC of processor so that I could not evaluate the theoretical peak performance. While I was doing program than faced problem with some JAVA time related functions and how to calculate cpu time and wall clock time.

## 2.2 Disk Benchmark

- a. AWS-EC2 Cloud had disk to measure a performance of it in terms of disk throughput(MB/s) and latency(ms). For Disk Benchmark I did Read and Write operation on file of 1B, 1KB, 1MB sequential and Random. Find result for multiple experiment like Sequential Read for 1B with 1-Thread, Random Write for 1KB with 2-thread and so on.

No of experiments =  $2 * 2 * 3 * 2 = 24$

(Read/Write) \* (Sequential/Random) \* (1B/1KB/1MB) \* (1-thread/2-thread)

Disk performance in terms of throughput(MB/s) and latency (ms) calculated for all the experiments mentioned below in table.

Practical throughput = **3745.318 MB/s**

Practical Latency = **0.267 ms**

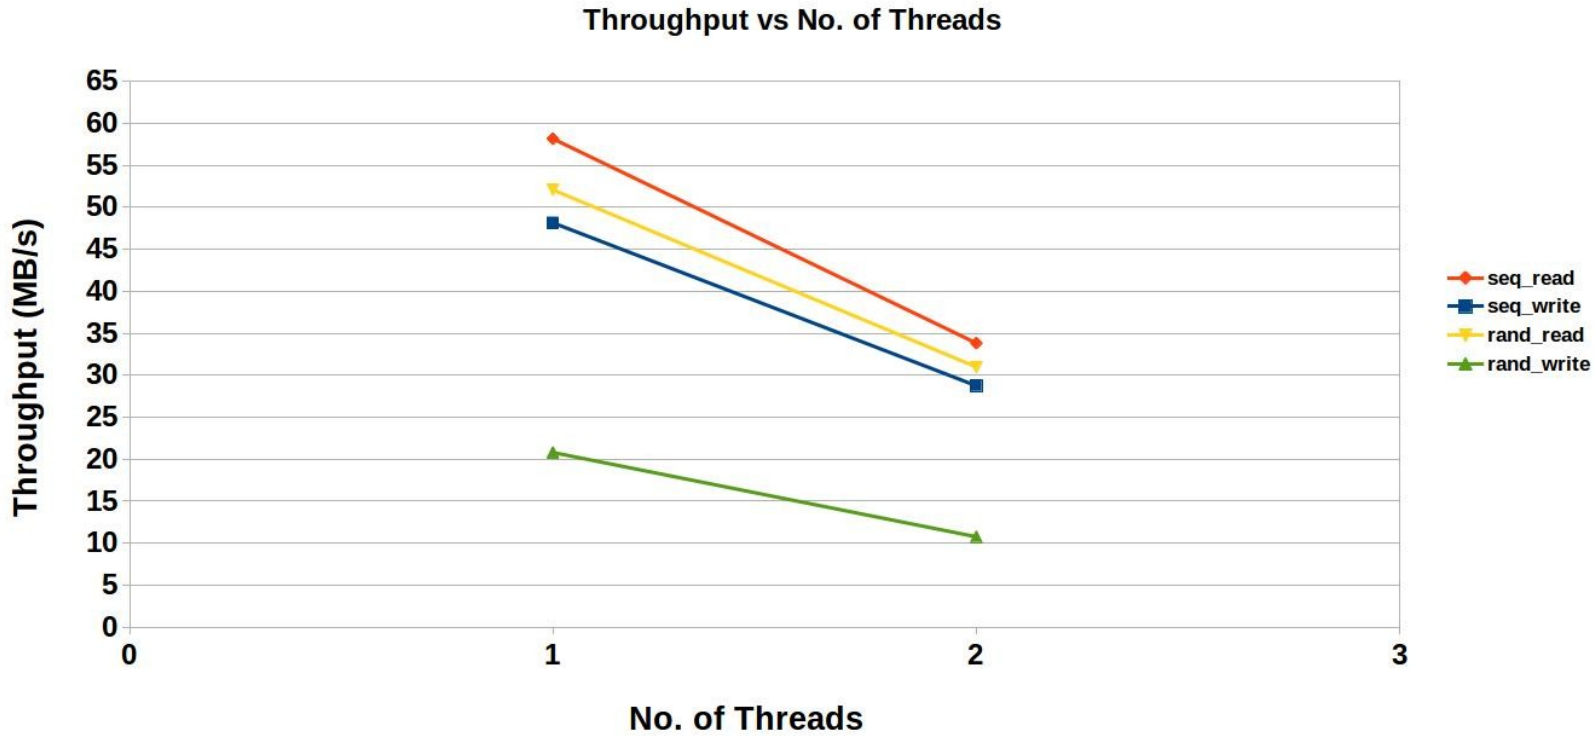
Theoretical Disk performance = 6.00 GB/s = 6000 MB/s (SATA V-3)

efficiency =  $(3754.318/6000) * 100 = 62.42 \%$

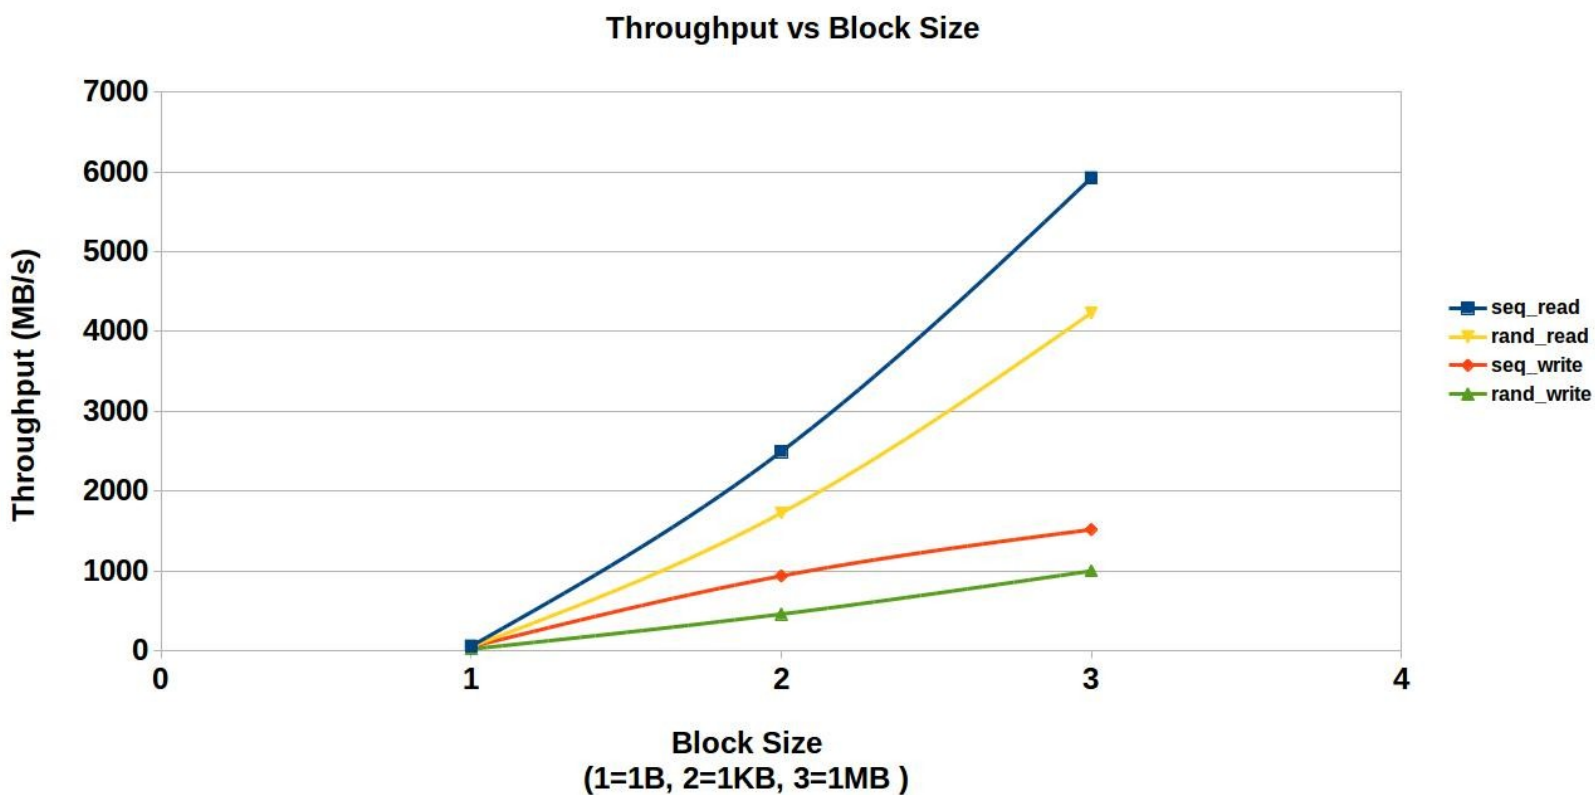
- Below graph is give the better idea about the comparison with the different experiments
- Sequential throughput always higher than Random throughput in both the case Read and Write.
- Write throughput is always less than the Read throughput and latency is reverse.
- Random access is costlier than Sequential access in terms of speed and accuracy.
- No of threads is increase than throughput is decrease in all cases because thread handling overhead.
- The lesser the latency, The Higher the performance.

- So, In terms of throughput(MB/s) and latency.

## Disk Throughput (1B Block)




## Disk Benchmark Throroghput (1 Thread)



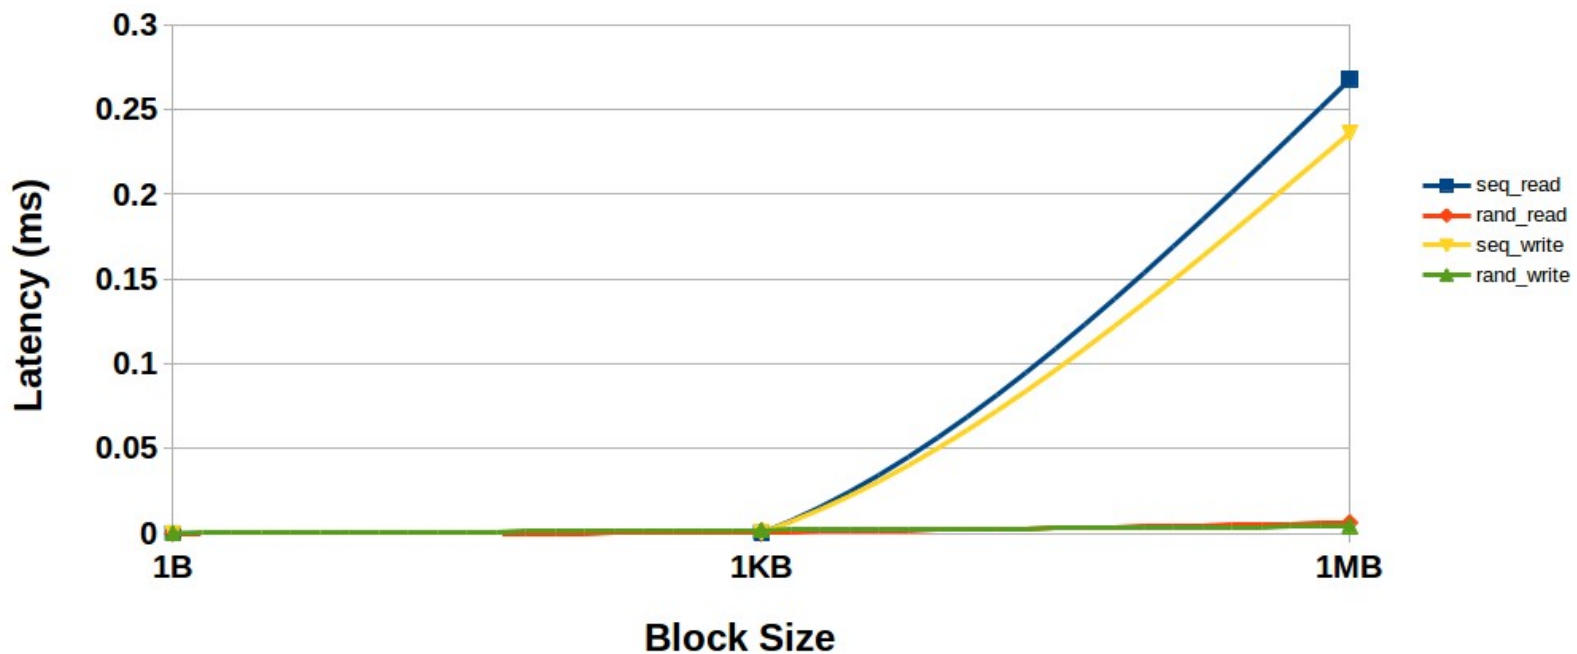
## Disk Latency (1B Block )

Latency (ms) vs No. of Threads



## Disk Latency

Latency (ms) vs Block Size





## Sequential Read

Block Size	No. of Threads	Average(Throughput (MB/s))	SD(Throughput (MB/s))	Average (Latency (ms))	SD(Latency(ms))
1B	1	56.5849	1.01223	0.000016	0.000004
1B	2	32.4231	0.96992	0.000030	0.000003
1KB	1	2484.4534	1.09461	0.000183	0.000040
1KB	2	2137.2445	0.94674	0.000190	0.000010
1MB	1	5880.31311	0.95105	0.267799	0.003500
1MB	2	5180.6477	1.0991	0.274545	0.002300

## Random Read

Block Size	No. of Threads	Average(Throughput (MB/s))	SD(Throughput (MB/s))	Average (Latency (ms))	SD(Latency(ms))
1B	1	52.071	0.64532	0.000019	0.000005
1B	2	30.9692	1.0403	0.000521	0.000032
1KB	1	1725.0571	0.67554	0.000798	0.000052
1KB	2	1566.4123	0.64284	0.000623	0.000076
1MB	1	4231.9685	0.51678	0.063230	0.000504
1MB	2	4023.3211	0.43572	0.167394	0.002034

## Sequential Write

Block Size	No. of Threads	Average(Throughput (MB/s))	SD(Throughput (MB/s))	Average (Latency (ms))	SD(Latency(ms))
1B	1	49.0506	0.6180	0.000020	0.000008
1B	2	30.8317	0.88293	0.000123	0.000042
1KB	1	935.9977	0.9232	0.000566	0.000048
1KB	2	690.2321	1.0022	0.000587	0.000072
1MB	1	1224.2568	1.2133	0.063200	0.004030
1MB	2	1023.3121	0.8992	0.083211	0.000342

## Random Write

Block Size	No. of Threads	Average(Throughput (MB/s))	SD(Throughput (MB/s))	Average (Latency (ms))	SD(Latency(ms))
1B	1	20.7958	0.9012	0.000069	0.000008
1B	2	10.7612	0.7812	0.000930	0.000078
1KB	1	508.9865	0.8233	0.001919	0.000204
1KB	2	412.3123	1.2110	0.001318	0.000123
1MB	1	998.5963	0.91212	0.042300	0.003279
1MB	2	732.1232	1.01311	0.303231	0.004321

## Note

- Average and Standard Deviation of throughput and latency gives the better and accurate performance of disk.
- Average and Standard Deviation are calculated for each and every experiments thrice.
- For Disk Performance in every experiments **File Size** is 100 MB.

## Practical Performance of Disk Speed on IOZONE Benchmark Tool:

- I used **iozone3\_434** for benchmarking of disk of AWS\_EC2 Cloud.
- Comparison Between Iozone result and My program Practical result for 1MB block size.

Access Type	Iozone	My Practical Results
Sequential Read	9485.232	5880.31311
Random Read	10230.845	4231.9685
Sequential Write	2192.645	1224.2568
Random Write	4282.950	998.5963

```
[ec2-user@ip-172-31-49-170 current]$ ./iozone -a -i 0 -i 1 -i 2 -s 1024
Iozone: Performance Test of File I/O
Version $Revision: 3.434 $
Compiled for 64 bit mode.
Build: linux

Contributors: William Norcott, Don Capps, Isom Crawford, Kirby Collins
              Al Slater, Scott Rhine, Mike Wisner, Ken Goss
              Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
              Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
              Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
              Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
              Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
              Vangel Bojaxhi, Ben England, Vikentsi Lapa,
              Alexey Skidanov.

Run began: Fri Feb 12 21:39:08 2016

Auto Mode
File size set to 1024 kB
Command line used: ./iozone -a -i 0 -i 1 -i 2 -s 1024
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.
```

	kB	reclen	write	rewrite	read	reread	random read	random write	bkwd read	record rewrite	stride read	fwrite	frewrite	fread	freread
1024	4	1735476	4146499	11249091	12456195	9485232	4282950								
1024	8	2443368	5169641	13472053	15516181	10454984	5226256								
1024	16	2516377	4995276	12944224	14044758	12944224	5042191								
1024	32	2716948	4995276	12037272	13998981	13102174	5885083								
1024	64	2349794	5169641	9569770	15027577	13686709	6093831								
1024	128	2192645	4451639	9485232	13998981	13264026	6059442								
1024	256	2250080	4195100	9402175	12173747	11249091	5821271								
1024	512	2254805	4614246	10039528	11903823	11249091	5630486								
1024	1024	2056183	4898425	11614118	10230845	10662628	5593820								

```
iozone test complete.
```

Fig : Iozone – 1MB multiple experiment

### Theoretical Performance of Disk

```
[ec2-user@ip-172-31-49-170 ~]$ sudo hdparm -tT /dev/sda

/dev/sda:
Timing cached reads:   21204 MB in  2.00 seconds = 10613.19 MB/sec
Timing buffered disk reads: 244 MB in  3.00 seconds = 81.24 MB/sec
[ec2-user@ip-172-31-49-170 ~]$
```

### EFFICIENCY:


$$\begin{aligned} \text{Efficiency of Iozone to Theoretical} &= (9485.232/10613.19) * 100 \\ &= 89.38 \% \end{aligned}$$

## 2.3 Memory Benchmark

- Memory benchmark is for Sequential and random Access for different data size of 1B/1KB/1MB.
- All different experiment is done thrice and took Average and Standard Deviation for better accuracy and consistency.
- Thread Effect on Memory Throughput
  - a.) As we increase the no of block size Sequential and Random throughput is increase
  - b.) As we increase the no of threads then Sequential and random throughput is decrease because thread overhead.
  - c.) For all experiment Sequential throughput is always higher than Random Throughput.


### Memory Throughput (MB/s) (1 - Thread)

Throughput ( MB) vs Block Size




## Throughput (MB/s) (1B Block)

Throughput (MB/s) vs No. of Threads




- Thread Effect on Memory Latency
  - a.) As we increase the no of block size Sequential and Random Latency is decrease
  - b.) As we increase the no of threads then Sequential and Random latency is increase because thread overhead.
  - c.) For all experiment Random Latency is always higher than Sequential Latency.

### Latency (ms) vs Block Size



### Latency (ms) vs No. of Threads



## Calculation

- Theoretical Throughput = **10600 MB/s** , Latency = **0.320 ms**
- The Optimal no of threads is 1, 1B Block size.
- Practical Throughput = **4710.211 MB/s** and Latency = **0.0000203 ms**
- Efficiency = **44.45 %**

## Sequential Access

Block Size	No. of Threads	Average(Throughput (MB/s))	SD(Throughput (MB/s))	Average (Latency (ms))	SD(Latency(ms))
1B	1	4710.211	1.01223	0.0000203	0.000004
1B	2	4444.312	0.96992	0.0000227	0.000003
1KB	1	4853.131	1.09461	0.0000194	0.000040
1KB	2	4564.502	0.94674	0.0000224	0.000010
1MB	1	5143.322	0.95105	0.0000176	0.003500
1MB	2	4670.34	1.0991	0.0000215	0.002300

## Random Access

Block Size	No. of Threads	Average(Throughput (MB/s))	SD(Throughput (MB/s))	Average (Latency (ms))	SD(Latency(ms))
1B	1	2698.231	0.64532	0.0000384	0.000005
1B	2	2444.98	1.0403	0.0000419	0.000032
1KB	1	2745.131	0.67554	0.0000378	0.000052
1KB	2	2493.05	0.64284	0.0000410	0.000076
1MB	1	2796.321	0.51678	0.0000370	0.000504
1MB	2	2570.56	0.43572	0.0000402	0.002034

## STREAM Benchmark Result Analysis

- I run Stream Benchmark and result is given in the below screen shot.
- Theoretical Throughput = **10600 MB/s** , Latency = **0.320 ms**
- While practical Throughput = **6714.45 MB/sec** and Latency = **0.03023 ms**
- Efficiency for Throughput = **63.35 %**
- Efficiency for Latency = **0.95 %**

```
[ec2-user@ip-172-31-49-170 Memory]$ gcc stream.c
[ec2-user@ip-172-31-49-170 Memory]$ ./a.out
-----
STREAM version $Revision: 5.10 $
-----
This system uses 8 bytes per array element.
-----
Array size = 100000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 28764 microseconds.
(= 28764 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:          5411.4      0.029769     0.029567     0.030050
Scale:         5315.3      0.030266     0.030102     0.030448
Add:           7699.2      0.031382     0.031172     0.031707
Triad:         7231.3      0.033374     0.033189     0.033571
-----
Solution Validates: avg error less than 1.000000e-13 on all three arrays
-----
[ec2-user@ip-172-31-49-170 Memory]$ █
```